


BT5110 Week 12 Lecture SQL, NoSQL & NewSQL

Chan Chee Yong
chancy@comp.nus.edu.sg

HOW TO WRITE A CV



Modern DBMSs: NoSQL & NewSQL



<http://www.informationweek.com/big-data/big-data-analytics/16-nosql-newsql-databases-to-watch/d/d-id/1269559>

Relational Database Systems



(Image: Software Engineering Daily)

Transactions

- Abstraction for representing a logical unit of work
- **ACID Properties**
 - ▶ **Atomicity**: Either all the effects of the transactions are reflected in the database or none are
 - ▶ **Consistency**: The execution of a transaction in isolation preserves the consistency of the database
 - ▶ **Isolation**: The execution of a transaction is isolated from the effects of other concurrent transaction executions
 - ▶ **Durability**: The effects of a committed transaction persists in the database even in the presence of system failures

Transaction Example: Money Transfer

```
1 int Transfer (int fromAcctId, int toAcctId, int amount)
2 {
3     EXEC SQL BEGIN DECLARE SECTION;
4         int fromBalance;    int toBalance;
5     EXEC SQL END DECLARE SECTION;
6     EXEC SQL WHENEVER SQLERROR GOTO query_error;
7
8     EXEC SQL SELECT balance INTO :fromBalance FROM Accounts
9         WHERE accountId = :fromAcctId;
10    if (fromBalance < amount) {
11        EXEC SQL ROLLBACK;    return 1;
12    }
13    EXEC SQL SELECT balance INTO :toBalance FROM Accounts
14        WHERE accountId = :toAcctId;
15    EXEC SQL UPDATE Accounts SET balance = :toBalance + :amount
16        WHERE accountId = :toAcctId;
17    EXEC SQL UPDATE Accounts SET balance = :fromBalance - :amount
18        WHERE accountId = :fromAcctId;
19    EXEC SQL COMMIT;
20    return 0;
21    query_error: printf ("SQL error: %ld\n", sqlca->sqlcode); exit();
22 }
```

Transaction Example: Money Transfer

Two possible execution outcomes:

Abort

```
begin transaction;  
  
select    balance into :fromBalance  
from      Accounts  
where     accountId = :fromAcctId;  
  
rollback;
```

Commit

```
begin transaction;  
  
select    balance into :fromBalance  
from      Accounts  
where     accountId = :fromAcctId;  
  
select    balance into :toBalance  
from      Accounts  
where     accountId = :toAcctId;  
  
update    Accounts  
set       balance = :toBalance + :amount  
where     accountId = 1;  
  
update    Accounts  
set       balance = :fromBalance - :amount  
where     accountId = 2;  
  
commit;
```

Serial Transaction Executions

Two possible serial executions of **Transfer(1,2,100)** & **Transfer(2,1,100)**

- (1): **begin transaction;**
select balance **into** :frombal **from** Accounts **where** accountId = 1;
select balance **into** :tobal **from** Accounts **where** accountId = 2;
update Accounts **set** balance = :tobal + 100 **where** accountId = 2;
update Accounts **set** balance = :frombal - 100 **where** accountId = 1;
commit;
begin transaction;
select balance **into** :frombal **from** Accounts **where** accountId = 2;
select balance **into** :tobal **from** Accounts **where** accountId = 1;
update Accounts **set** balance = :tobal + 100 **where** accountId = 1;
update Accounts **set** balance = :frombal - 100 **where** accountId = 2;
commit;
- (2): **begin transaction;**
select balance **into** :frombal **from** Accounts **where** accountId = 2;
select balance **into** :tobal **from** Accounts **where** accountId = 1;
update Accounts **set** balance = :tobal + 100 **where** accountId = 1;
update Accounts **set** balance = :frombal - 100 **where** accountId = 2;
commit;
begin transaction;
select balance **into** :frombal **from** Accounts **where** accountId = 1;
select balance **into** :tobal **from** Accounts **where** accountId = 2;
update Accounts **set** balance = :tobal + 100 **where** accountId = 2;
update Accounts **set** balance = :frombal - 100 **where** accountId = 1;
commit;

Concurrent Transaction Executions

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Serializable Transaction Executions

- A concurrent execution of a set of transactions is **serializable** if this execution is equivalent to some serial execution
- A concurrent execution CE is **equivalent** to a serial execution SE if
 - ▶ both CE & SE produce the same final database state, and
 - ▶ every read operation in CE returns the same value as the corresponding read operation in SE
- A serial execution is trivially serializable

Example 1: Non-serializable Execution

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal:
tobal:

Accounts

(1, \$400)
(2, \$2000)

Transfer(2,1,100)

frombal:
tobal:

Example 1: Non-serializable Execution

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal:

Accounts

(1, \$400)

(2, \$2000)

Transfer(2,1,100)

frombal:

tobal:

Example 1: Non-serializable Execution

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal: 2000

Accounts

(1, \$400)

(2, \$2000)

Transfer(2,1,100)

frombal:

tobal:

Example 1: Non-serializable Execution

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400
tobal: 2000

Accounts

(1, \$400)
(2, \$2100)

Transfer(2,1,100)

frombal:
tobal:

Example 1: Non-serializable Execution

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400
tobal: 2000

Accounts

(1, \$400)
(2, \$2100)

Transfer(2,1,100)

frombal: 2100
tobal:

Example 1: Non-serializable Execution

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400
tobal: 2000

Accounts

(1, \$400)
(2, \$2100)

Transfer(2,1,100)

frombal: 2100
tobal: 400

Example 1: Non-serializable Execution

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

Transfer(1,2,100)

frombal: 400
tobal: 2000

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

Accounts

(1, \$300)
(2, \$2100)

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(2,1,100)

frombal: 2100
tobal: 400

Example 1: Non-serializable Execution

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

Transfer(1,2,100)

frombal: 400
tobal: 2000

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

Accounts

(1, \$500)
(2, \$2100)

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(2,1,100)

frombal: 2100
tobal: 400

Example 1: Non-serializable Execution

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

Transfer(1,2,100)

frombal: 400
tobal: 2000

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

Accounts

(1, \$500)
(2, \$2000)

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(2,1,100)

frombal: 2100
tobal: 400

Example 1: Non-serializable Execution

A concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal: 2000

Accounts

(1, \$500)

(2, \$2000)

Transfer(2,1,100)

frombal: 2100

tobal: 400

Example 2: Serializable Execution

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal:

tobal:

Accounts

(1, \$400)

(2, \$2000)

Transfer(2,1,100)

frombal:

tobal:

Example 2: Serializable Execution

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal:

Accounts

(1, \$400)

(2, \$2000)

Transfer(2,1,100)

frombal:

tobal:

Example 2: Serializable Execution

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal: 2000

Accounts

(1, \$400)

(2, \$2000)

Transfer(2,1,100)

frombal:

tobal:

Example 2: Serializable Execution

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal: 2000

Accounts

(1, \$400)

(2, \$2100)

Transfer(2,1,100)

frombal:

tobal:

Example 2: Serializable Execution

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal: 2000

Accounts

(1, \$400)

(2, \$2100)

Transfer(2,1,100)

frombal: 2100

tobal:

Example 2: Serializable Execution

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal: 2000

Accounts

(1, \$300)

(2, \$2100)

Transfer(2,1,100)

frombal: 2100

tobal:

Example 2: Serializable Execution

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal: 2000

Accounts

(1, \$300)

(2, \$2100)

Transfer(2,1,100)

frombal: 2100

tobal: 300

Example 2: Serializable Execution

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal: 2000

Accounts

(1, \$400)

(2, \$2100)

Transfer(2,1,100)

frombal: 2100

tobal: 300

Example 2: Serializable Execution

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal: 2000

Accounts

(1, \$400)

(2, \$2000)

Transfer(2,1,100)

frombal: 2100

tobal: 300

Example 2: Serializable Execution

Another concurrent execution of **Transfer(1,2,100)** & **Transfer(2,1,100)**

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 1;

select balance **into** :tobal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :tobal + 100 **where** accountId = 2;

begin transaction;

select balance **into** :frombal **from** Accounts **where** accountId = 2;

update Accounts **set** balance = :frombal - 100 **where** accountId = 1;

commit;

select balance **into** :tobal **from** Accounts **where** accountId = 1;

update Accounts **set** balance = :tobal + 100 **where** accountId = 1;

update Accounts **set** balance = :frombal - 100 **where** accountId = 2;

commit;

Transfer(1,2,100)

frombal: 400

tobal: 2000

Accounts

(1, \$400)

(2, \$2000)

Transfer(2,1,100)

frombal: 2100

tobal: 300

SQL Isolation Levels

- The **isolation level** for a transaction affects what the transaction will read
- SQL defines four isolation levels
 - ▶ Read Uncommitted (**weakest isolation level**)
 - ▶ Read Committed
 - ▶ Repeatable Read
 - ▶ Serializable (**strongest isolation level**)
- Choice of isolation level affects correctness vs performance tradeoff
- In many DBMSs, the default isolation level is **Read Committed**
- Configure using **set transaction isolation level** statement

Creating Indexes to Speed up Queries

- Consider the following SQL query:

```
SELECT *  
FROM Customers  
WHERE state = 'CA' AND city = 'Santa Barbara'
```

- How to create an index on attributes (state,city) of Customers table?

```
CREATE INDEX ON Customers (state,city)
```

Creating Indexes to Speed up Queries (cont.)

- Index on (state,city) could speed up the following queries:
 - `SELECT * FROM Customers WHERE state = 'CA' AND city = 'Napa'`
 - `SELECT * FROM Customers WHERE state = 'CA' AND city < 'Napa'`
 - `SELECT * FROM Customers WHERE state = 'CA' AND city > 'Napa'`
 - `SELECT * FROM Customers WHERE state >= 'CA'`
- However, the index is not useful for these queries:
 - `SELECT * FROM Customers WHERE city = 'Springfield'`
 - `SELECT * FROM Customers WHERE state > 'CA' AND city = 'Napa'`

Scalar Subqueries

Sales

name	month	amount
Alice	2018-01-01	400
Alice	2018-02-01	300
Alice	2018-03-01	600
Bob	2018-01-01	200
Bob	2018-02-01	250
Bob	2018-03-01	300
Carol	2018-01-01	250
Carol	2018-02-01	200
Carol	2018-03-01	400

name	month	amount	total
Alice	2018-01-01	400	1300
Alice	2018-02-01	300	1300
Alice	2018-03-01	600	1300
Bob	2018-01-01	200	750
Bob	2018-02-01	250	750
Bob	2018-03-01	300	750
Carol	2018-01-01	250	850
Carol	2018-02-01	200	850
Carol	2018-03-01	400	850

```
SELECT      *, (SELECT SUM(S2.amount)
                  FROM Sales S2
                  WHERE S2.name = S.name)
FROM        Sales S
ORDER BY   name, month
```

Common Table Expressions (CTEs)

Sales

name	month	amount
Alice	2018-01-01	400
Alice	2018-02-01	300
Alice	2018-03-01	600
Bob	2018-01-01	200
Bob	2018-02-01	250
Bob	2018-03-01	300
Carol	2018-01-01	250
Carol	2018-02-01	200
Carol	2018-03-01	400

name	month	amount	total
Alice	2018-01-01	400	1300
Alice	2018-02-01	300	1300
Alice	2018-03-01	600	1300
Bob	2018-01-01	200	750
Bob	2018-02-01	250	750
Bob	2018-03-01	300	750
Carol	2018-01-01	250	850
Carol	2018-02-01	200	850
Carol	2018-03-01	400	850

WITH TotalSales AS

(SELECT name, SUM(amount) AS total

FROM Sales

GROUP BY name)

SELECT S.*, T.total

FROM Sales S NATURAL JOIN TotalSales T

ORDER BY S.name, S.month

Window Functions

Sales

name	month	amount
Alice	2018-01-01	400
Alice	2018-02-01	300
Alice	2018-03-01	600
Bob	2018-01-01	200
Bob	2018-02-01	250
Bob	2018-03-01	300
Carol	2018-01-01	250
Carol	2018-02-01	200
Carol	2018-03-01	400

name	month	amount	total
Alice	2018-01-01	400	1300
Alice	2018-02-01	300	1300
Alice	2018-03-01	600	1300
Bob	2018-01-01	200	750
Bob	2018-02-01	250	750
Bob	2018-03-01	300	750
Carol	2018-01-01	250	850
Carol	2018-02-01	200	850
Carol	2018-03-01	400	850

```
SELECT *, SUM(amount) OVER W AS total
FROM Sales
WINDOW W AS (PARTITION BY name)
ORDER BY name, month
```

Window Functions: Running Total

Sales

name	month	amount
Alice	2018-01-01	400
Alice	2018-02-01	300
Alice	2018-03-01	600
Bob	2018-01-01	200
Bob	2018-02-01	250
Bob	2018-03-01	300
Carol	2018-01-01	250
Carol	2018-02-01	200
Carol	2018-03-01	400

name	month	amount	runningTotal
Alice	2018-01-01	400	400
Alice	2018-02-01	300	700
Alice	2018-03-01	600	1300
Bob	2018-01-01	200	200
Bob	2018-02-01	250	450
Bob	2018-03-01	300	750
Carol	2018-01-01	250	250
Carol	2018-02-01	200	450
Carol	2018-03-01	400	850

```
SELECT *, SUM(amount) OVER W AS runningtotal
FROM Sales
WINDOW W AS (PARTITION BY name ORDER BY month)
ORDER BY name, month
```

Window Functions: Revenue Growth

Sales

name	month	amount
Alice	2018-01-01	400
Alice	2018-02-01	300
Alice	2018-03-01	600
Bob	2018-01-01	200
Bob	2018-02-01	250
Bob	2018-03-01	300
Carol	2018-01-01	250
Carol	2018-02-01	200
Carol	2018-03-01	400

month	revenue	prevMthRevenue	revenueGrowth
2018-01-01	850	null	null
2018-02-01	750	850	-11.76
2018-03-01	1300	750	73.33

- Revenue = total monthly sales
- Revenue Growth = $\frac{R - R'}{R'} \times 100$
 - ▶ R = revenue for a given month
 - ▶ R' = revenue for preceding month

Window Functions: Revenue Growth (cont.)

Sales

name	month	amount
Alice	2018-01-01	400
Alice	2018-02-01	300
Alice	2018-03-01	600
Bob	2018-01-01	200
Bob	2018-02-01	250
Bob	2018-03-01	300
Carol	2018-01-01	250
Carol	2018-02-01	200
Carol	2018-03-01	400

month	revenue	prevMthRevenue	revenueGrowth
2018-01-01	850	null	null
2018-02-01	750	850	-11.76
2018-03-01	1300	750	73.33

WITH MthlySales AS

(SELECT month, SUM(amount) AS revenue
FROM Sales GROUP BY month),

RevenuePair AS (

SELECT month, revenue,
lag(revenue) OVER W AS prevMthRevenue


FROM MthlySales
WINDOW W AS (ORDER BY month))

SELECT *, ROUND(100.0 * (revenue-prevMthRevenue)/prevMthRevenue,2)
AS revenueGrowth


FROM RevenuePair
ORDER BY month

NoSQL Database Systems

- Supports **large-scale data management** challenges of today's web-based applications
 - ▶ Database Scalability, High Availability, Low Latency
 - ▶ Schema-less data or data with dynamic schema
- Modern **distributed database systems**
 - ▶ Data being sharded & replicated across a cluster of servers



Data Sharding (Image: Oracle)



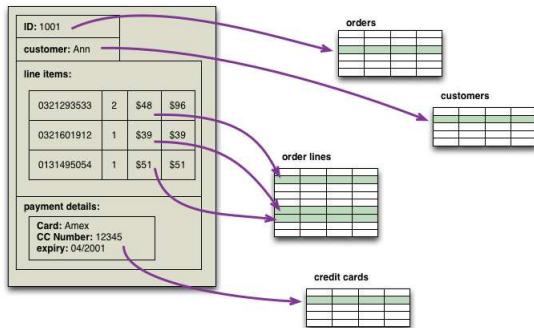
Data Replication (Image: Lloyd, et. al, SOSP 2011)

Early NoSQL Database Systems

- Schema-less data
 - ▶ Key-value store
- Simple access API instead of query language
 - ▶ Put, Get, Delete
- Limited/No ACID transactional support
- Weak consistency for replicated data
 - ▶ Eventual consistency

NoSQL Database Systems

- **Key-value stores** (e.g., Dynamo, Redis)
- **Column-family stores** (e.g., BigTable, Cassandra, HBase)
- **Document stores** (e.g., Cosmos DB, MarkLogic, MongoDB)



(Martin Fowler, 2012)


- **Graph database systems** (e.g., JanusGraph, Neo4j)

Graph Database Systems

- Based on different **graph data models**: property graph, RDF graph, hypergraph
- **Property Graphs**
 - ▶ Systems: JanusGraph, Neo4j, etc.
- **RDF Graphs**
 - ▶ RDF = Resource Description Framework
 - ▶ Data stores known as *triplestores* / semantic graph databases
 - ★ Store data as (subject, predicate, object) triples
 - ▶ Query language: SPARQL
 - ▶ Supports RDF Schema (RDFS) & Web Ontology Language (OWL) inference
 - ▶ Uses: Linked Open Data, Knowledge Graphs, etc.
 - ▶ Systems: AllegroGraph, GraphDB, etc.
- **Hypergraphs**
 - ▶ Systems: HyperGraphDB, Microsoft Graph Engine, etc.

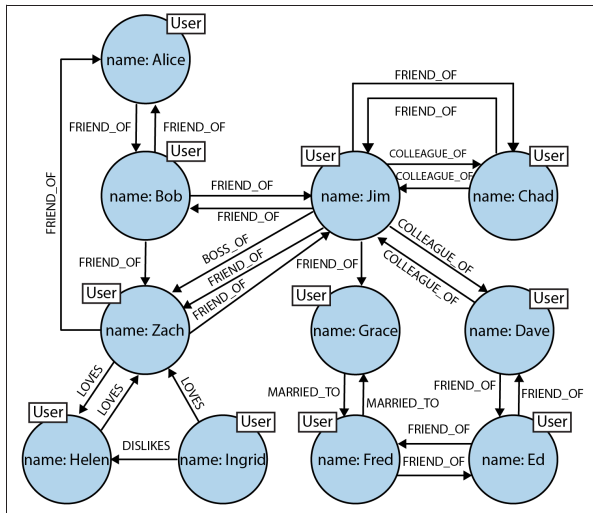
Property Graph Data Model

- **Nodes** represent entities
 - ▶ Each node has at least one label & possibly properties
- Directed edges represent **relationships** between entities
 - ▶ Each relationship has a type & possibly properties
- Each **property** is a key-value pair



(<http://neo4j-contrib.github.io/developer-resources/get-started/graph-database>)

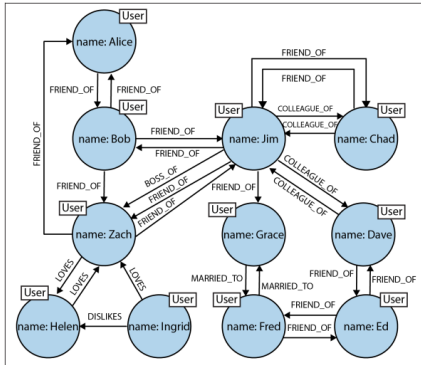
Property Graph Data Model (cont.)



(Robinson, Webber & Eifrem, 2015)

Graph Query Example

Find all users who are friends of Bob that share similar friends as Bob



Bob's friends: Alice, Jim, Zach

Alice's friends: Bob

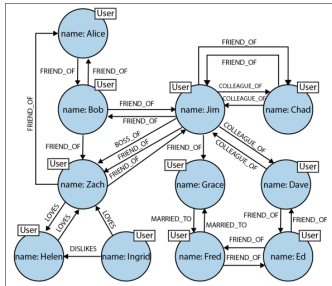
Jim's friends: Bob, Chad, Grace, Zach

Zach's friends: Alice, Jim

Neo4j's Cypher Query Language

- Declarative query language based on property graph model

```
MATCH (e)-[:FRIEND_OF]-(bob)-[:FRIEND_OF]->(f)-[:FRIEND_OF]->(e)
WHERE bob.name = "Bob"
RETURN f.name AS name,
        count(e) AS score,
        collect(e.name) AS friends
ORDER BY score DESC
```

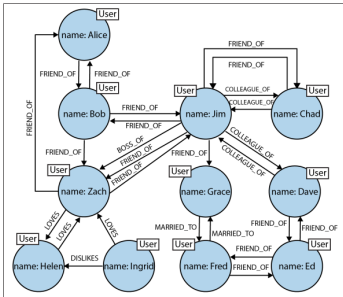


name	score	friends
"Zach"	2	["Alice", "Jim"]
"Jim"	1	["Zach"]

Another Cypher Query

For each user, find the number of his/her direct/indirect friends

```
MATCH (u:User)  
OPTIONAL MATCH (u)-[:FRIEND_OF*]->(u2:User)  
RETURN u.name AS name,  
count(DISTINCT u2) AS numFriends  
ORDER BY name
```



name	numFriends
"Alice"	5
"Bob"	5
"Chad"	5
"Dave"	2
"Ed"	2
"Fred"	2
"Grace"	0
"Helen"	0
"Ingrid"	0
"Jim"	5
"Zach"	5

NewSQL Database Systems

- Targeted at OLTP workloads
- Features
 - ▶ Relational data model
 - ▶ SQL query language
 - ▶ ACID transactions
 - ▶ Runs on distributed cluster of shared-nothing nodes
- Some examples:
 - ▶ Clustrix
 - ▶ CockroachDB
 - ▶ Google Spanner
 - ▶ MemSQL
 - ▶ VoltDB

Database-as-a-Service (DBaaS)

- **RDBMS**

- ▶ **Amazon RDS** (Amazon, Aurora, MySQL, MariaDB, SQL Server, Oracle, PostgreSQL)
<https://aws.amazon.com/rds/>
- ▶ **Google Cloud SQL** (MySQL, PostgreSQL)
<https://cloud.google.com/sql/>

- **NoSQL**

- ▶ **Amazon DynamoDB**
<https://aws.amazon.com/dynamodb/>
- ▶ **Microsoft Azure Cosmos DB**
<https://azure.microsoft.com/en-us/services/cosmos-db/>

- **NewSQL**

- ▶ **Google Cloud Spanner**
<https://cloud.google.com/spanner/>